

Jingxuan Ma

+1 (410) 240-4366 | jma107@jh.edu | [LinkedIn](#) | [Portfolio](#)

Data Science M.S. student at Johns Hopkins with a background in Applied Statistics and data analytics. Skilled in R, advanced SQL, Python ETL pipelines using Apache Airflow, Tableau dashboard development, operations automation, stakeholder collaboration, and AI prompt engineering to accelerate data analysis.

EDUCATION

Johns Hopkins University August 2025 - June 2027
Master of Science in Data Science Baltimore, US

• **Coursework:** Database Systems | Introduction to Algorithms | Nonlinear Optimization | Human-Computer Interaction

University College Cork September 2021 - June 2025
Bachelor of Science in Actuarial Sciences Cork, Ireland | Beijing, China

• **Coursework:** Probability & Statistics | Actuarial Mathematics | Financial Mathematics | Regression Analysis | Statistical Computing

PROFESSIONAL EXPERIENCE

AI Agent Intern | Beijing Yiling Network Technology Co., Ltd. Beijing, China | June 2026 - Present

- Built a Python agent that projects risk-analyst resumes from a locked master by selecting credit/claims/Monte Carlo evidence and injecting content-only OOXML updates
- Designed format-lock and quality-gate checks (shell fingerprint, hyperlink preservation, Word PDF one-page validation, evidence-linked bullets) so delivery copies keep master fonts, margins, and list styles without rebuilding the document
- Used fixture JDs spanning risk and pricing roles to validate that tailored bullets stay evidence-linked and the delivery PDF remains exactly one page

Data Analyst Intern | Shenwan Hongyuan Securities Co., Ltd. Beijing, China | June 2024 - August 2024

- Faced with the need to evaluate whether a pricing model library could meet both runtime and maintainability requirements, conducted a structured feasibility analysis across Python, pure & optimized C++, Eigen vectorization, OpenMP, and ctypes-based Python/C++ integration
- Built and benchmarked a 100,000-path Monte Carlo pricing simulation, applying compiler optimization, vectorized matrix operations, and multithreaded random-path generation to isolate major bottlenecks in computation-intensive pricing workflows
- Delivered a hybrid architecture recommendation that assigned heavy simulation and statistical computation to C++ while keeping configuration, preprocessing, and visualization in Python; reduced runtime from approx. 33.4s to approx. 1.4s with OpenMP optimization

Data Analyst Intern | Yinhua Fund Management Co., Ltd. Beijing, China | June 2023 - August 2023

- Given the need to support campaign review and market reporting, collected, organized, and cleaned market and customer data from multiple business materials, creating structured datasets for trend analysis and management communication
- Used Python, Excel, and visualization techniques to analyze market patterns, prepare charts and analytical materials, and convert raw performance data into insights that could support investment-product marketing decisions
- Prepared concise reporting assets for internal stakeholders by translating quantitative findings into business narratives, improving clarity around customer behavior, campaign performance, and market trends

PROJECTS

Insurance Claims Severity Modeling | Python, R, pandas, scikit-learn, Monte Carlo Simulation Independent Project

- To estimate claim severity across heterogeneous policy segments, collected and cleaned historical claims data and engineered exposure- and policy-level features to support distributional analysis of loss costs
- Analyzed skewed loss distributions, outliers, and cost drivers; compared regression and tree-based approaches to evaluate trade-offs among predictive accuracy, robustness, and business interpretability for underwriting and reserving use cases
- Extended the modeling framework with Monte Carlo simulation and performance-benchmarking concepts to evaluate scalable pricing workflows, translating model outputs into risk segmentation, pricing review, and portfolio-level reporting insights

Tesla Vehicle Quality & Risk Analytics Pipeline | Python, Apache Airflow, SQL, Tableau, NHTSA API Independent Project

- Built an end-to-end ETL pipeline extracting 11,349 NHTSA complaints and 60 recall actions across Tesla Models 3/S/X/Y (2015–2024); engineered dynamic model-name discovery and ID-based deduplication to reduce 30,777 raw records to 11,349 verified distinct complaints
- Engineered frequency- and severity-based risk indicators (crash/fire/injury-weighted scores by model-year and component) in SQL, applying an actuarial-style frequency $\times$ severity loss framework to quantify and rank vehicle quality risk across product lines
- Built and published an interactive Tableau dashboard visualizing component-level risk rankings and recall-lag trends (first complaint to official recall), translating 11,349 complaint records into decision-ready risk insights for quality and safety stakeholders.

• COMPETITIONS

Mathematical Contest in Modeling | Team Leader Remote | February 2023

- Led a student modeling team through problem scoping, data cleaning, indicator screening, visualization, model development, and technical report writing, coordinating responsibilities to deliver a structured solution under time constraints

Mathematical Modeling for College Students | Team Member Remote | September 2022

- Collected and processed data for statistical modeling tasks; built models using Python and R, analyzed results with SPSS, and contributed to result interpretation and final written recommendations

SKILLS & CERTIFICATIONS

Python, R, SQL, Tableau, data cleaning, feature engineering, scikit-learn, XGBoost, Credit Risk, Apache Airflow, exploratory analysis, Pandas, NumPy, Monte Carlo Simulation, Claims Modeling, stakeholder reporting